

Project ML

```
import numpy as np
import pandas as pd

import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.model_selection import train_test_split

import tensorflow as tf

from sklearn.metrics import confusion_matrix, classification_report


data = pd.read_csv('../content/emotions.csv')

data

{"type": "dataframe", "variable_name": "data"}

sample = data.loc[0, 'fft_0_b': 'fft_749_b']

plt.figure(figsize=(16, 10))
plt.plot(range(len(sample)), sample)
plt.title("Features fft_0_b through fft_749_b")
plt.show()
```



```
data['label'].value_counts()
```

```
label
NEUTRAL    716
NEGATIVE   708
POSITIVE   708
Name: count, dtype: int64
```

```
label_mapping = {'NEGATIVE': 0, 'NEUTRAL': 1, 'POSITIVE': 2}
```

Preprocessing

```
def preprocess_inputs(df):
    df = df.copy()

    df['label'] = df['label'].replace(label_mapping)

    y = df['label'].copy()
    X = df.drop('label', axis=1).copy()

    X_train, X_test, y_train, y_test = train_test_split(X, y,
train_size=0.7, random_state=123)

    return X_train, X_test, y_train, y_test
```

```
X_train, X_test, y_train, y_test = preprocess_inputs(data)

<ipython-input-63-3f2c0931f06d>:4: FutureWarning: Downcasting behavior
in `replace` is deprecated and will be removed in a future version. To
retain the old behavior, explicitly call
`result.infer_objects(copy=False)`. To opt-in to the future behavior,
set `pd.set_option('future.no_silent_downcasting', True)`
  df['label'] = df['label'].replace(label_mapping)

X_train

{"type": "dataframe", "variable_name": "X_train"}
```

Modeling

```
inputs = tf.keras.Input(shape=(X_train.shape[1],))

# Wrap tf.expand_dims in a Lambda layer
expand_dims = tf.keras.layers.Lambda(lambda x: tf.expand_dims(x,
axis=2))(inputs)

gru = tf.keras.layers.GRU(256, return_sequences=True)(expand_dims)

flatten = tf.keras.layers.Flatten()(gru)

outputs = tf.keras.layers.Dense(3, activation='softmax')(flatten)

model = tf.keras.Model(inputs=inputs, outputs=outputs)
print(model.summary())

Model: "functional_4"
```

Layer (type) Param #	Output Shape
input_layer_8 (InputLayer) 0	(None, 2548)
lambda_4 (Lambda) 0	(None, 2548, 1)
gru_4 (GRU) 198,912	(None, 2548, 256)

0	flatten_4 (Flatten)	(None, 652288)
1,956,867	dense_4 (Dense)	(None, 3)

Total params: 2,155,779 (8.22 MB)

Trainable params: 2,155,779 (8.22 MB)

Non-trainable params: 0 (0.00 B)

None

```
model.compile(
    optimizer='adam',
    loss='sparse_categorical_crossentropy',
    metrics=['accuracy']
)
```

```
history = model.fit(
    X_train,
    y_train,
    validation_split=0.2,
    batch_size=32,
    epochs=50,
    callbacks=[
        tf.keras.callbacks.EarlyStopping(
            monitor='val_loss',
            patience=5,
            restore_best_weights=True
        )
    ]
)
```

Epoch 1/50

38/38 ————— 6s 125ms/step - accuracy: 0.6867 - loss: 41.4652 - val_accuracy: 0.8930 - val_loss: 14.5967

Epoch 2/50

38/38 ————— 5s 124ms/step - accuracy: 0.9349 - loss: 3.0479 - val_accuracy: 0.9365 - val_loss: 3.9214

Epoch 3/50

38/38 ————— 5s 121ms/step - accuracy: 0.9339 - loss: 3.1856 - val_accuracy: 0.8629 - val_loss: 8.7536

Epoch 4/50

38/38 ————— 5s 125ms/step - accuracy: 0.9459 - loss:

```

2.5055 - val_accuracy: 0.9465 - val_loss: 4.1846
Epoch 5/50
38/38 ━━━━━━━━━━━ 5s 118ms/step - accuracy: 0.9873 - loss:
0.2460 - val_accuracy: 0.9431 - val_loss: 3.9431
Epoch 6/50
38/38 ━━━━━━━━━━━ 4s 117ms/step - accuracy: 0.9864 - loss:
0.3796 - val_accuracy: 0.9732 - val_loss: 1.2569
Epoch 7/50
38/38 ━━━━━━━━━━━ 5s 120ms/step - accuracy: 0.9836 - loss:
0.2925 - val_accuracy: 0.9498 - val_loss: 4.3082
Epoch 8/50
38/38 ━━━━━━━━━━━ 5s 116ms/step - accuracy: 0.9707 - loss:
2.1224 - val_accuracy: 0.9565 - val_loss: 1.6326
Epoch 9/50
38/38 ━━━━━━━━━━━ 5s 118ms/step - accuracy: 0.9955 - loss:
0.1921 - val_accuracy: 0.9398 - val_loss: 5.3997
Epoch 10/50
38/38 ━━━━━━━━━━━ 4s 117ms/step - accuracy: 0.9818 - loss:
2.0692 - val_accuracy: 0.9465 - val_loss: 6.0800
Epoch 11/50
38/38 ━━━━━━━━━━━ 5s 114ms/step - accuracy: 0.9887 - loss:
0.2426 - val_accuracy: 0.9431 - val_loss: 4.5341

```

Results

```

model_acc = model.evaluate(X_test, y_test, verbose=0)[1]
print("Test Accuracy: {:.3f}%".format(model_acc * 100))

Test Accuracy: 96.406%


y_pred = np.array(list(map(lambda x: np.argmax(x),
model.predict(X_test))))

cm = confusion_matrix(y_test, y_pred)
clr = classification_report(y_test, y_pred,
target_names=label_mapping.keys())

plt.figure(figsize=(8, 8))
sns.heatmap(cm, annot=True, vmin=0, fmt='g', cbar=False, cmap='Blues')
plt.xticks(np.arange(3) + 0.5, label_mapping.keys())
plt.yticks(np.arange(3) + 0.5, label_mapping.keys())
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.title("Confusion Matrix")
plt.show()

print("Classification Report:\n-----\n", clr)

```



Classification Report:

	precision	recall	f1-score	support
NEGATIVE	0.95	0.96	0.95	201
NEUTRAL	1.00	0.98	0.99	231
POSITIVE	0.94	0.95	0.95	208
accuracy			0.96	640

macro avg	0.96	0.96	0.96	640
weighted avg	0.96	0.96	0.96	640